



Introduction to Deep Learning

Amrita Vishwa Vidyapeetham
Amritapuri Campus






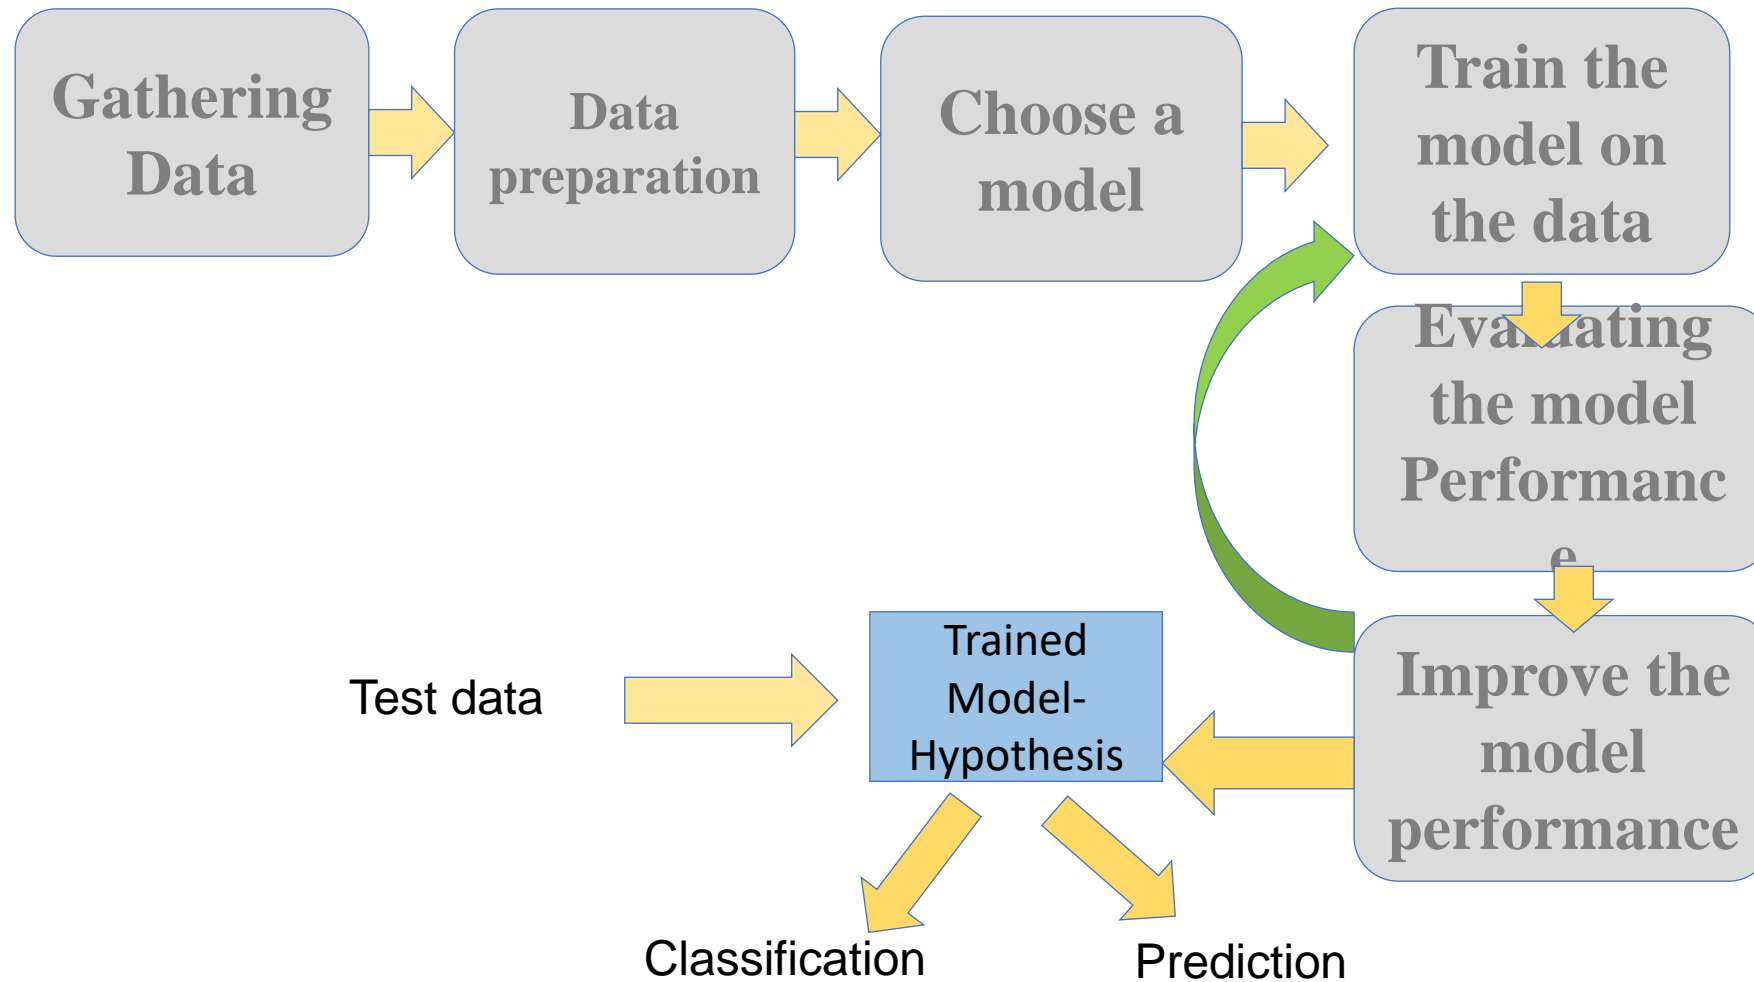
Introduction to Deep Learning

Citation Note: The concept, the content, and the structure of this article were inspired by the awesome lectures and the material offered by Prof. [Mitesh M. Khapra](#) on [NPTEL's Deep Learning](#) course

Ingredients- ML model



Recap :Supervised Learning



Recap

Id	Cl.thickness	Cell.size	Cell.shape	Marg.adhesion	Epith.c.size	Bare.nuclei	Bl.cromatin	Normal.nucleoli	Mitoses	Class
1000025	5	1	1	1	2	1	3	1	1	benign
1002945	5	4	4	5	7	10	3	2	1	benign
1015425	3	1	1	1	2	2	3	1	1	benign
1016277	6	8	8	1	3	4	3	7	1	benign
1017023	4	1	1	3	2	1	3	1	1	benign
1017122	8	10	10	8	7	10	9	7	1	malignant
1018099	1	1	1	1	2	10	3	1	1	benign
1018561	2	1	2	1	2	1	3	1	1	benign
1033078	2	1	1	1	2	1	1	1	5	benign
1033078	4	2	1	1	2	1	2	1	1	benign
1035283	1	1	1	1	1	1	3	1	1	benign
1036172	2	1	1	1	2	1	2	1	1	benign
1041801	5	3		3	2	3	4	4	1	malignant

Train the model with the above data set

Is my cancer Malignant or benign

NO it is benign

Test the model with a new real time input.

1041801	5	3	3	3	2	3	4	4	1	
---------	---	---	---	---	---	---	---	---	---	--



VectorStock

Biological Neuron



Takes an input, processes it, throws out an output.

Dendrite: Receives signals from other neurons

Soma: Processes the information

Axon: Transmits the output of this neuron

Synapse: Point of connection to other neurons




Courtesy: Towardsdatascience.com

Biological Neuron

Each neuron gets fired/activated only when its respective criteria is met

Our sense organs interact with the outer world and send the visual and sound information to the neurons

. Let's say you are watching a video clip. Now the information your brain receives is taken in by the **“laugh or not”** set of neurons that will help you make a decision on whether to laugh or not.




There is a massively parallel interconnected network of 10^{11} neurons (100 billion) in our brain and their connections are not simple

Courtesy: Towardsdatascience.com


Neurons are arranged in a hierarchical fashion

Neurons are arranged in a hierarchical fashion and each layer has its own role and responsibility.

To detect a face, the brain could be relying on the entire network and not on a single layer.



Layer 1: detect edges & corners



Layer 2: form feature groups



Layer 3: detect high level objects, faces, etc.

Courtesy :Sample illustration of hierarchical processing. Credits: Mitesh M. Khapra's lecture slides



The very first step towards the perceptron we use today was taken in 1943 by McCulloch and Pitts, by mimicking the functionality of a biological neuron.

Recall Biological Neuron

Artificial Neuron

The fundamental Building block of Deep Learning

The most fundamental unit of deep neural networks is called an *artificial neuron/perceptron*





Courtesy :: Mitesh M. Khapra's lecture slides

Artificial Neuron

“threshold logic” to mimic the thought process

The fundamental Building block of Deep Learning

Artificial Neural Networks



McCulloch-Pitts Neuron

The first computational model of a neuron was proposed by Warren McCulloch (neuroscientist) and Walter Pitts (logician) in 1943

It may be divided into 2 parts. The first part, g takes an input (ahem dendrite ahem), performs an aggregation and based on the aggregated value the second part, f makes a decision.

Lets suppose that I want to **predict my own decision**, **whether to watch a random football game or not on TV?**.

The inputs are all boolean i.e., $\{0, 1\}$ and my output variable is also boolean $\{0, 1\}$: Will watch it, 1: Won't watch it}.


So, x_1 could be *isPremierLeagueOn* (I like Premier League more)

• x_2 could be *isItAFriendlyGame* (I tend to care less about the friendlies)

• x_3 could be *isNotHome* (Can't watch it when I'm running errands. Can I?)

• x_4 could be *isManUnitedPlaying* (I am a big Man United fan. GGMU!) and so on.

Excitatory or inhibitory inputs.



These inputs can either be *excitatory* or *inhibitory*.

Inhibitory inputs are those that have maximum effect on the decision making irrespective of other inputs

i.e., if x_3 is 1 (not home) then my output will always be 0 i.e., the neuron will never fire, so x_3 is an inhibitory input

Excitatory inputs are NOT the ones that will make the neuron fire on their own but they might fire it when combined together


$$g(x_1, x_2, x_3, \dots, x_n) = g(\mathbf{x}) = \sum_{i=1}^n x_i$$

$$y = f(g(\mathbf{x})) = \begin{cases} 1 & \text{if } g(\mathbf{x}) \geq \theta \\ 0 & \text{if } g(\mathbf{x}) < \theta \end{cases}$$

The mathematical Model- McCulloch-Pitts Neuron (MP Neuron)

- ◆ McCulloch and Pitts proposed a highly simplified computational model of the neuron.
- ◆ g aggregates the inputs and the function f takes a decision based on this aggregation.
- ◆ The inputs can be excitatory or inhibitory

These inputs can either be *excitatory* or *inhibitory*




$y = 0$ if any x_i is inhibitory, else

$$g(x_1, x_2, \dots, x_n) = g(x) = \sum_{i=1}^n x_i$$

$$y = f(g(x)) = 1 \text{ if } g(x) \geq b$$

$$= 0 \text{ if } g(x) < b$$

What task can be done with MP Neuron



Pitch in line	Impact	Missing stumps	Is it LBW? (y)
1	0	0	0
0	1	1	0
1	1	1	1
0	1	0	




Courtesy :: Mitesh M. Khapra's lecture slides

Loss Function

We aim at minimizing the loss function

									
Launch (within 6 months)	0	1	1	0	0	1	0	1	0
Weight (<160g)	1	0	1	0	0	0	1	0	0
Screen size (<5.9 in)	1	0	1	0	1	0	1	0	1
dual sim	1	1	0	0	0	1	0	1	0
Internal memory (>= 64 GB, 4GB RAM)	1	1	1	1	1	1	1	1	0
NFC	0	1	1	0	1	0	1	1	1
Radio	1	0	0	1	1	1	0	0	0
Battery(>3500mAh)	0	0	0	1	0	1	0	1	0
Price > 20k	0	1	1	0	0	0	1	1	1
Like? (y)	1	0	1	0	1	1	0	1	0
prediction \hat{y}	1	0	0	1	1	1	1	0	0
loss	0	0	1	-1	0	0	-1	1	0



$$\hat{y} = \sum_{i=1}^n x_i \geq b$$

$$\hat{y} = x_1 + x_2 \geq b$$

$$loss = \sum_i (y_i - \hat{y}_i)^2$$

$$loss = \sum_i y_i - \hat{y}_i$$

Courtesy :: Mitesh M. Khapra's lecture slides

Training data

Launch (within 6 months)	0	1	1	0	0	1	0	1	1	0
Weight (<160g)	1	0	1	0	0	0	1	0	0	1
Screen size (<5.9 In)	1	0	1	0	1	0	1	0	1	0
dual sim	1	1	0	0	0	1	0	1	0	0
Internal memory (>= 64 GB, 4GB RAM)	1	1	1	1	1	1	1	1	1	0
NFC	0	1	1	0	1	0	1	1	1	0
Radio	1	0	0	1	1	1	0	0	0	0
Battery(>3500mAh)	0	0	0	1	0	1	0	1	0	0
Price > 20k	0	1	1	0	0	0	1	1	1	0
Like? (y)	1	1	1	0	0	1	1	1	0	0
predicted	1	1	0	1	1	1	1	0	0	0

Test data

1	0	0	1
0	1	1	1
0	1	1	1
0	1	0	0
1	0	0	0
0	0	1	0
1	1	1	0
1	1	1	0
0	0	1	0
0	1	0	0
0	1	1	0



Learning


MODEL

Predictions

$$\text{Accuracy} = \frac{\text{Number of correct predictions}}{\text{Total number of predictions}}$$


$$= \frac{3}{4} = 75\%$$

Boolean Functions Using M-P Neuron




sum \geq theta, the neuron will fire otherwise, it won't.

AND Function




An AND function neuron would only fire when ALL the inputs are ON
i.e., **$g(\mathbf{x}) \geq 3$** here.

OR Function




neuron would fire if ANY of the inputs is ON
i.e., **$g(\mathbf{x}) \geq 1$** here.

Geometric Interpretation OR function



OR function


$$x_1 + x_2 = \sum_{i=1}^2 x_i \geq 1$$



$x_1 + x_2 = 1$ to graphically show that all those inputs whose output when passed through the OR function M-P neuron lie ON or ABOVE that line and all the input points that lie BELOW that line are going to output 0.


The M-P neuron just learnt a linear decision boundary!

OR Function With 3 Inputs




$$x_1 + x_2 + x_3 = \sum_{i=1}^3 x_i \geq 1$$



Lets just generalize this by looking at a 3 input OR function M-P unit. In this case, the possible inputs are 8 points — $(0,0,0)$, $(0,0,1)$, $(0,1,0)$, $(1,0,0)$, $(1,0,1)$,... you got the point(s). We can map these on a 3D graph and this time we draw a decision boundary in 3 dimensions.



The plane that satisfies the decision boundary equation $x_1 + x_2 + x_3 = 1$ is shown below:




Boolean Functions Using M-P Neuron



In this case, the decision boundary equation is $x_1 + x_2 = 2$. Here, all the input points that lie ON or ABOVE, just $(1,1)$, output 1 when passed through the AND function M-P neuron. It fits! The decision boundary works!

Boolean Functions Using M-P Neuron




$$x_1 \text{ AND } !x_2^*$$

inhibitory input i.e., x_2 so whenever x_2 is 1, the output will be 0

$x_1 \text{ AND } !x_2$ would output 1 only when x_1 is 1 and x_2 is 0 so it is obvious that the threshold parameter should be 1.

$$g(x) \text{ i.e., } x_1 + x_2 \text{ would be } \geq 1$$

NOR Function



For a NOR neuron to fire, we want ALL the inputs to be 0 so the thresholding parameter should also be 0 and we take them all as inhibitory input.

NOT Function



For a NOT neuron, 1 outputs 0 and 0 outputs 1. So we take the input as an inhibitory input and set the thresholding parameter to 0. It works!

Limitations Of M-P Neuron

- What about non-boolean (say, real) inputs?
- Do we always need to hand code the threshold?
- Are all inputs equal? What if we want to assign more importance to some inputs?
- What about functions which are not linearly separable? Say XOR function.



Frank Rosenblatt

Overcoming the limitations of the M-P neuron, **Frank Rosenblatt**, an American psychologist, proposed the classical **perception model**, the mighty *artificial neuron*, in 1958. It is more generalized computational model than the McCulloch-Pitts neuron where weights and thresholds can be learnt over time

Namah Shivaya