

```
In [1]: ▶ print("""
Immadi Shyam Prasad    (AM.EN.U4CSE19164)
D B R S Praneeth Varma (AM.EN.U4CSE19265)
D S K Phani Chyavan    (AM.EN.U4CSE19163)
""")
```

```
Immadi Shyam Prasad    (AM.EN.U4CSE19164)
D B R S Praneeth Varma (AM.EN.U4CSE19265)
D S K Phani Chyavan    (AM.EN.U4CSE19163)
```

```
In [1]: ▶ import torch
import torchvision
from torchvision import transforms, datasets
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.autograd import Variable
from torch.utils.data import DataLoader
from torch.utils.data import sampler
from torch.utils.data import random_split
from torch.utils.data import SubsetRandomSampler
from random import sample

import matplotlib.pyplot as plt

import glob
from PIL import Image as I

import pandas as pd

import os

import numpy as np

import timeit
```

```
In [2]: ▶ pwd
```

```
Out[2]: '/mnt/batch/tasks/shared/LS_root/mounts/clusters/mmahadevareddy1/code/Users/mmahadevareddy/cvision/project1'
```

In [3]:

```
#https://www.kaggle.com/c/state-farm-distracted-driver-detection/data

#path of test,train,validation
train_dataset_path = "/mnt/batch/tasks/shared/LS_root/mounts/clusters/mmahadev"
test_dataset_path = "/mnt/batch/tasks/shared/LS_root/mounts/clusters/mmahadev"
val_dataset_path = "/mnt/batch/tasks/shared/LS_root/mounts/clusters/mmahadev"

mean = [0.3124, 0.3782, 0.3708] # found by calculation net.batch_mean_and_std
std = [0.2778, 0.3213, 0.3222]

#data argumentaion resizing,normalizing,converting to tensor

train_transforms = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(torch.Tensor(mean),
        torch.Tensor(std))
])

test_transforms = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(torch.Tensor(mean),
        torch.Tensor(std))
])

val_transforms = transforms.Compose([
    transforms.Resize((224,224)),
    transforms.ToTensor(),
    transforms.Normalize(torch.Tensor(mean),
        torch.Tensor(std))
])

train_dataset = datasets.ImageFolder(root=train_dataset_path,transform=train_transforms)
test_dataset = datasets.ImageFolder(root=test_dataset_path,transform=test_transforms)
val_dataset = datasets.ImageFolder(root=val_dataset_path,transform=val_transforms)

#data Loading
train_loader = torch.utils.data.DataLoader(dataset = train_dataset,batch_size=10)
test_loader = torch.utils.data.DataLoader(dataset = test_dataset,batch_size=10)
val_loader = torch.utils.data.DataLoader(dataset = val_dataset,batch_size=10)
```

```
In [4]: ▶ #function for printing images in grid
def show_transformed_images(dataset):
    loader = torch.utils.data.DataLoader(dataset, batch_size=6, shuffle=True)
    for data in loader:
        images, labels = data
        break
    print(labels)
    print(images.view(-1, 224*224).shape)
    grid = torchvision.utils.make_grid(images, nrow=3)
    plt.figure(figsize=(11, 11))
    plt.imshow(np.transpose(grid, (1, 2, 0)))
    plt.show()

show_transformed_images(train_dataset)
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

```
tensor([1, 6, 9, 6, 0, 5])
torch.Size([18, 50176])
```



```
In [5]: ▶ torch.cuda.is_available()
```

```
Out[5]: True
```

```
In [6]: ▶ def my_plot(epochs, loss,title1,x_label1,y_label1):
    plt.plot(epochs, loss)
    plt.title(title1)
    plt.xlabel(x_label1)
    plt.ylabel(y_label1)
    return plt.show()
def just_plot(accu,accu1,title1,y_label1):
    epochs = range(1,4)
    plt.plot(epochs, accu, 'g', label='Training accuracy')
    plt.plot(epochs, accu1, 'b', label='validation accuracy')
    plt.title(title1)
    plt.xlabel('Epochs')
    plt.ylabel(y_label1)
    plt.legend()
    return plt.show()
```

```
In [7]: ▶
num_classes = 10
net = torchvision.models.alexnet()
net.fc = nn.Linear(100, num_classes)
gpu_net = net.to('cuda')

import torch.optim as optim
from sklearn.metrics import accuracy_score

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(gpu_net.parameters(), lr=0.001, momentum=0.9, nesterov=
```

In [8]: ▶

```
train_loss= []
train_accuracy=[]

num_epochs = 3
for epoch in range(num_epochs): # loop over the dataset multiple times
    epoch_loss= []
    running_train_loss = 0.0
    total=0
    correct = 0
    print('train load')
    for i, data in enumerate(train_loader):
        # get the inputs; data is a list of [inputs, labels]
        inputs, labels = data
        inputs = inputs.to('cuda')
        labels = labels.to('cuda')

        # zero the parameter gradients
        optimizer.zero_grad()

        # forward + backward + optimize
        outputs = gpu_net(inputs)

        loss = criterion(outputs, labels)
        epoch_loss.append(loss.item())

        loss.backward()
        optimizer.step()

        for idx, k in enumerate(outputs):
            if torch.argmax(k) == labels[idx]:
                correct += 1
            total += 1

        # print statistics
        running_train_loss += loss.item()
        if i % 50 == 49: # print every 50 mini-batches
            print('[%d, %5d] Train loss: %.3f' %
                  (epoch + 1, i + 1, running_train_loss / 50))
            running_train_loss = 0.0

    train_accuracy.append(round(correct/total,3))
    train_loss.append(sum(epoch_loss)/len(epoch_loss))
    running_val_loss = 0.0

print('Finished Training')
```

```
[1, 1200] Train loss: 1.400
train load
[2, 50] Train loss: 1.273
[2, 100] Train loss: 1.289
```

```
[2, 100] Train loss: 1.205
[2, 150] Train loss: 1.207
[2, 200] Train loss: 1.220
[2, 250] Train loss: 1.226
[2, 300] Train loss: 1.022
[2, 350] Train loss: 0.918
[2, 400] Train loss: 0.987
[2, 450] Train loss: 0.996
[2, 500] Train loss: 0.982
[2, 550] Train loss: 0.944
[2, 600] Train loss: 0.829
[2, 650] Train loss: 0.823
[2, 700] Train loss: 0.743
[2, 750] Train loss: 0.725
[2, 800] Train loss: 0.606
[2, 850] Train loss: 0.734
[2, 900] Train loss: 0.701
```

```

In [9]: ▶ val_loss= []
          val_accuracy=[]

          num_epochs = 3
          for epoch in range(num_epochs):
              print('val load')
              epoch_loss= []
              running_val_loss = 0.0
              correct = 0
              total=0
              for i, data in enumerate(val_loader):
                  # get the inputs; data is a list of [inputs, labels]
                  inputs, labels = data
                  inputs = inputs.to('cuda')
                  labels = labels.to('cuda')

                  # zero the parameter gradients
                  optimizer.zero_grad()

                  # forward + backward + optimize
                  outputs = gpu_net(inputs)

                  loss = criterion(outputs, labels)
                  epoch_loss.append(loss.item())
                  loss.backward()
                  optimizer.step()
                  for idx, j in enumerate(outputs):
                      if torch.argmax(j) == labels[idx]:
                          correct += 1
                      total += 1

                  running_train_loss += loss.item()
                  if i % 50 == 49:      # print every 50 mini-batches
                      print('[%d, %5d] Train loss: %.3f' %
                            (epoch + 1, i + 1, running_train_loss / 50))
                      running_train_loss = 0.0

              val_accuracy.append(round(correct/total,3))
              val_loss.append(sum(epoch_loss)/len(epoch_loss))
              running_val_loss = 0.0
          print('validation')

[1, 100] Train loss: 0.301
[1, 150] Train loss: 0.279
[1, 200] Train loss: 0.326
[1, 250] Train loss: 0.265
[1, 300] Train loss: 0.339
val load
[2, 50] Train loss: 0.448
[2, 100] Train loss: 0.150
[2, 150] Train loss: 0.250
[2, 200] Train loss: 0.225
[2, 250] Train loss: 0.197
[2, 300] Train loss: 0.208
val load
[3, 50] Train loss: 0.425

```

```
[3, 100] Train loss: 0.190
[3, 150] Train loss: 0.138
[3, 200] Train loss: 0.150
[3, 250] Train loss: 0.192
[3, 300] Train loss: 0.192
validation
```


```
In [11]: ▶ print(train_accuracy, val_accuracy)
```

```
[0.232, 0.692, 0.864] [0.906, 0.935, 0.946]
```


```
In [12]: ▶ print(train_loss, val_loss)
```

```
[2.385819746556337, 0.8672663599325865, 0.41273387137468165] [0.29770814210
98237, 0.21012411083836632, 0.17033227367585466]
```

```
In [13]: ▶ just_plot(train_accuracy, val_accuracy, "train and validation accuracy", "accuracy")
```



```
In [14]: ▶ just_plot(train_loss, val_loss, "train and validation loss", "loss")
```





```
In [15]: ▶ #testing model
correct = 0
total = 0
with torch.no_grad():
    for data in test_loader:
        x, y = data
        output = gpu_net(x.to('cuda'))
        for idx, i in enumerate(output):
            if torch.argmax(i) == y[idx]:
                correct += 1
            total += 1
print('accuracy:\t',round(correct/total,3))
```

accuracy: 0.937

```
In [16]: ▶ x, y = next(iter(test_loader))
```

```
In [17]: ▶ x[0]
```

```
Out[17]: tensor([[[[-0.8987, -0.8987, -0.8987, ..., 2.1928, 2.2069, 1.8964],
  [-0.8987, -0.8987, -0.8987, ..., 2.0940, 2.2493, 2.2493],
  [-0.8987, -0.8987, -0.8987, ..., 1.8540, 2.0517, 2.1505],
  ...,
  [-0.7293, -0.7575, -0.7999, ..., -0.9693, -0.9834, -0.9834],
  [-0.7716, -0.7999, -0.8422, ..., -0.9834, -0.9834, -0.9834],
  [-0.8563, -0.8846, -0.9128, ..., -0.9834, -0.9834, -0.9834]],

  [[-0.9696, -0.9696, -0.9696, ..., 1.9231, 1.8742, 1.5813],
  [-0.9696, -0.9696, -0.9696, ..., 1.8620, 1.9231, 1.9109],
  [-0.9696, -0.9696, -0.9696, ..., 1.7034, 1.8376, 1.8986],
  ...,
  [-0.7621, -0.7865, -0.8231, ..., -1.0672, -1.0795, -1.0672],
  [-0.7987, -0.8231, -0.8598, ..., -1.0795, -1.0795, -1.0672],
  [-0.8720, -0.8964, -0.9208, ..., -1.0795, -1.0795, -1.0672]],

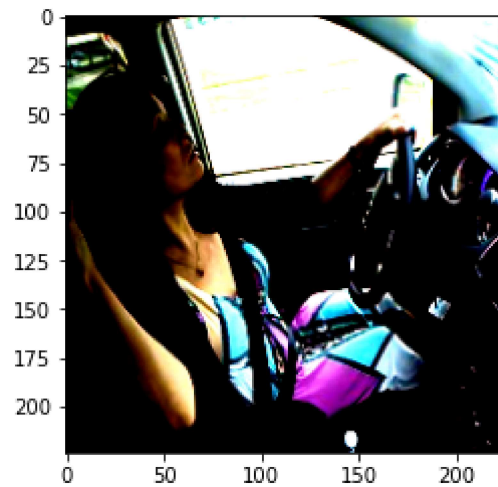
  [[-1.0048, -1.0048, -1.0048, ..., 1.6729, 1.6972, 1.4538],
  [-1.0048, -1.0048, -1.0048, ..., 1.5999, 1.7581, 1.7703],
  [-1.0048, -1.0048, -1.0048, ..., 1.4295, 1.6120, 1.6972],
  ...,
  [-0.7370, -0.7614, -0.7979, ..., -1.0291, -1.0535, -1.0656],
  [-0.7735, -0.7979, -0.8344, ..., -1.0413, -1.0535, -1.0656],
  [-0.8466, -0.8709, -0.8952, ..., -1.0413, -1.0535, -1.0656]]]])
```

In [18]: `#checking for single images`

```
print(y[8])  
plt.imshow(np.transpose(x[8],(1,2,0)))  
plt.show()
```

Clipping input data to the valid range for imshow with RGB data ([0..1] for floats or [0..255] for integers).

tensor(8)



```

In [19]: ► output = gpu_net(x.to('cuda'))
          print(torch.max(output.data,1))
          for idx, i in enumerate(output):
              if torch.argmax(i) == y[idx]:
                  print(' correct',torch.argmax(i),y[idx])
              else:
                  print('incorrect',torch.argmax(i),y[idx])

          torch.return_types.max(
          values=tensor([17.5003, 28.5840, 41.8555, 18.6444, 28.8190, 28.0235, 20.824
          3, 21.7966,
                        25.3960, 29.9994], device='cuda:0'),
          indices=tensor([3, 4, 6, 9, 3, 3, 7, 7, 8, 7], device='cuda:0'))
          uncorrect tensor(3, device='cuda:0') tensor(0)
              correct tensor(4, device='cuda:0') tensor(4)
              correct tensor(6, device='cuda:0') tensor(6)
              correct tensor(9, device='cuda:0') tensor(9)
              correct tensor(3, device='cuda:0') tensor(3)
              correct tensor(3, device='cuda:0') tensor(3)
              correct tensor(7, device='cuda:0') tensor(7)
              correct tensor(7, device='cuda:0') tensor(7)
              correct tensor(8, device='cuda:0') tensor(8)
              correct tensor(7, device='cuda:0') tensor(7)

```

In [ ]: ►

In [ ]: ►