



Introduction to Deep Learning

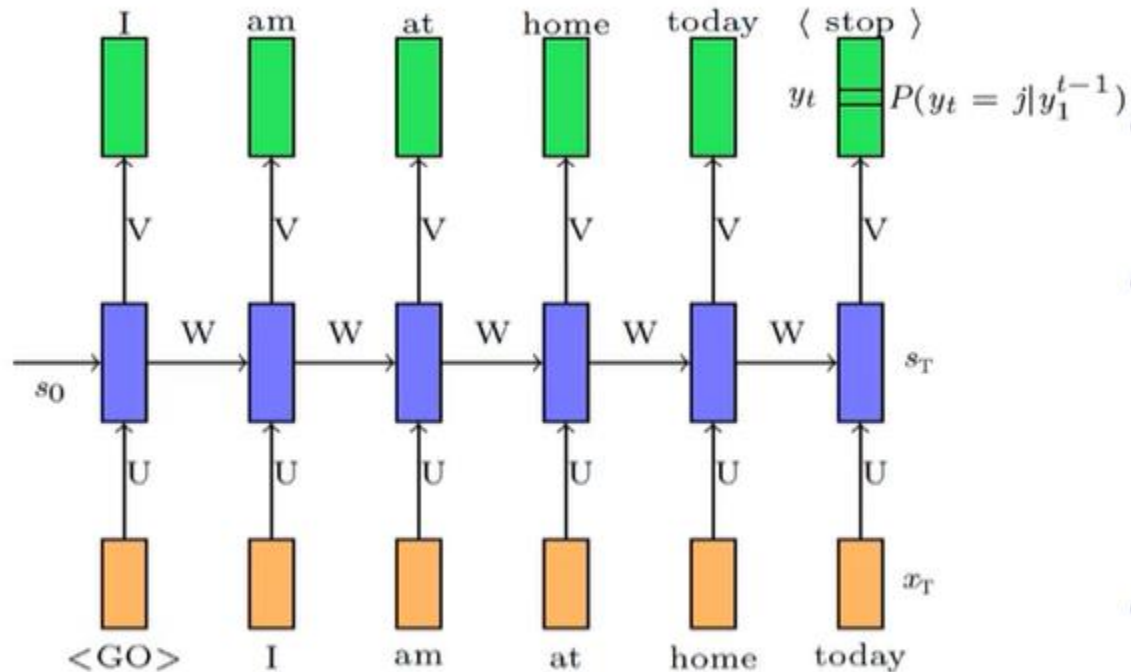
Amrita Vishwa Vidyapeetham
Amritapuri Campus





Encoder- Decoder Models

Recap - RNN

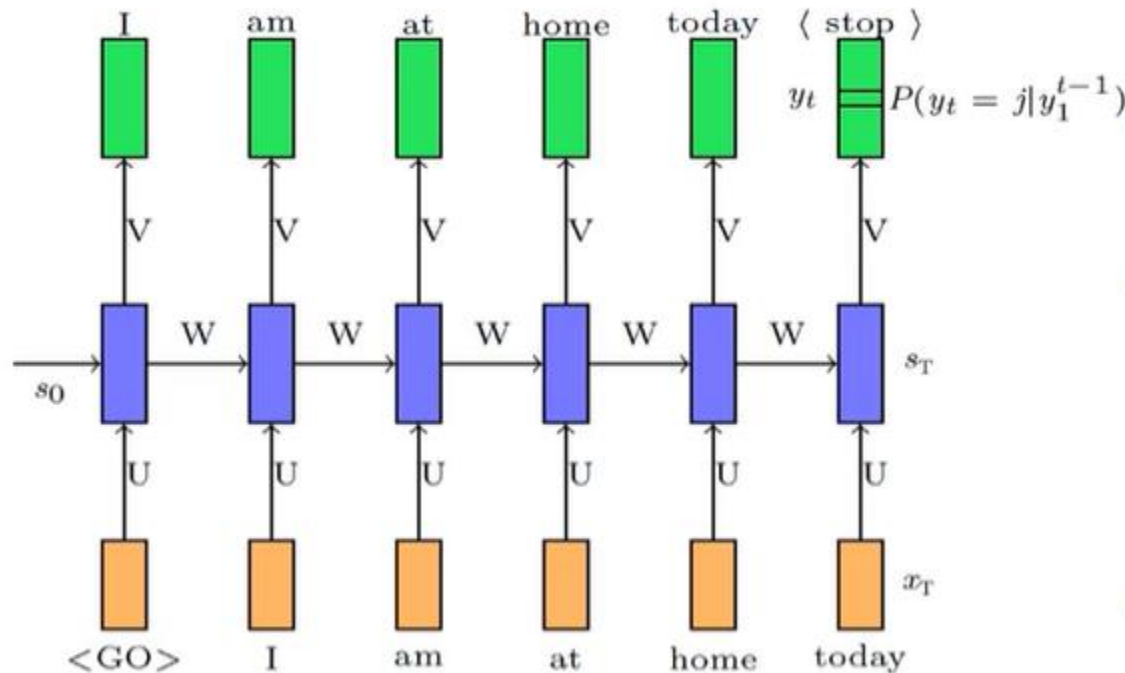


- We will start by revisiting the problem of language modeling
- Informally, given ' $t - i$ ' words we are interested in predicting the t^{th} word
- More formally, given y_1, y_2, \dots, y_{t-1} we want to find

$$y^* = \operatorname{argmax} P(y_t | y_1, y_2, \dots, y_{t-1})$$

- Let us see how we model $P(y_t | y_1, y_2, \dots, y_{t-1})$ using a RNN
- We will refer to $P(y_t | y_1, y_2, \dots, y_{t-1})$ by shorthand notation: $P(y_t | y_1^{t-1})$

Recap - RNN



- We are interested in

$$P(y_t = j | y_1, y_2 \dots y_{t-1})$$

where $j \in V$ and V is the set of all vocabulary words

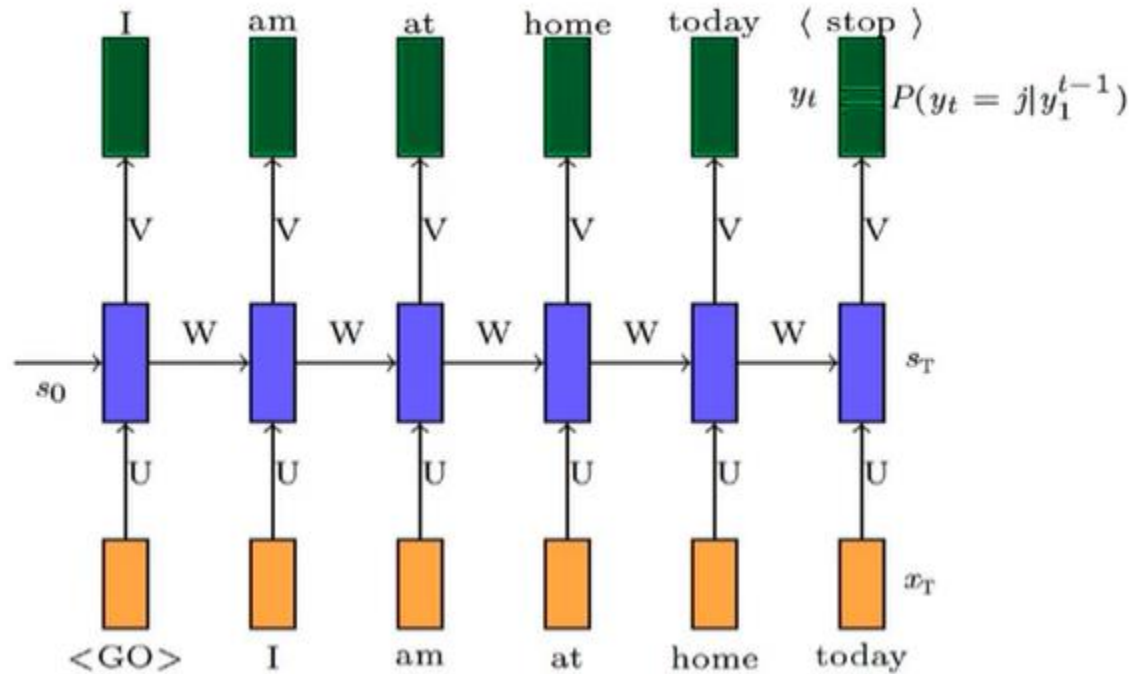
- Using an RNN we compute this as

$$P(y_t = j | y_1^{t-1}) = \text{softmax}(Vs_t + c)_j$$

- In other words we compute

$$\begin{aligned} P(y_t = j | y_1^{t-1}) &= P(y_t = j | s_t) \\ &= \text{softmax}(Vs_t + c)_j \end{aligned}$$

Recap - RNN



Data:

India, officially the Republic of India, is a country in South Asia. It is the seventh-largest country by area,

- **Data:** All sentences from any large corpus (say wikipedia)

- **Model:**

$$s_t = \sigma(Ws_{t-1} + Ux_t + b)$$

$$P(y_t = j | y_1^{t-1}) = \text{softmax}(Vs_t + c)_j$$

- **Parameters:** U, V, W, b, c

- **Loss:**

$$\mathcal{L}(\theta) = \sum_{t=1}^T \mathcal{L}_t(\theta)$$

$$\mathcal{L}_t(\theta) = -\log P(y_t = \ell_t | y_1^{t-1})$$

$$s_t = \sigma(U x_t + W s_{t-1} + b) \quad \tilde{s}_t = \sigma(W(o_t \odot s_{t-1}) + U x_t + b)$$

$$s_t = i_t \odot s_{t-1} + (1 - i_t) \odot \tilde{s}_t$$

$$s_t = \text{RNN}(s_{t-1}, x_t)$$

$$s_t = \text{GRU}(s_{t-1}, x_t)$$

$$\tilde{s}_t = \sigma(W h_{t-1} + U x_t + b)$$

$$s_t = f_t \odot s_{t-1} + i_t \odot \tilde{s}_t$$

$$h_t = o_t \odot \sigma(s_t)$$

$$h_t, s_t = \text{LSTM}(h_{t-1}, s_{t-1}, x_t)$$

Encoder – Decoder for Image Captioning

Image Captioning Problem

- So far we have seen how to model the conditional probability distribution $P(y_t|y_1^{t-1})$
- More informally, we have seen how to generate a sentence given previous words
- What if we want to generate a sentence given an image?

- Earlier we modeled $P(y_t|y_1^{t-1})$ as

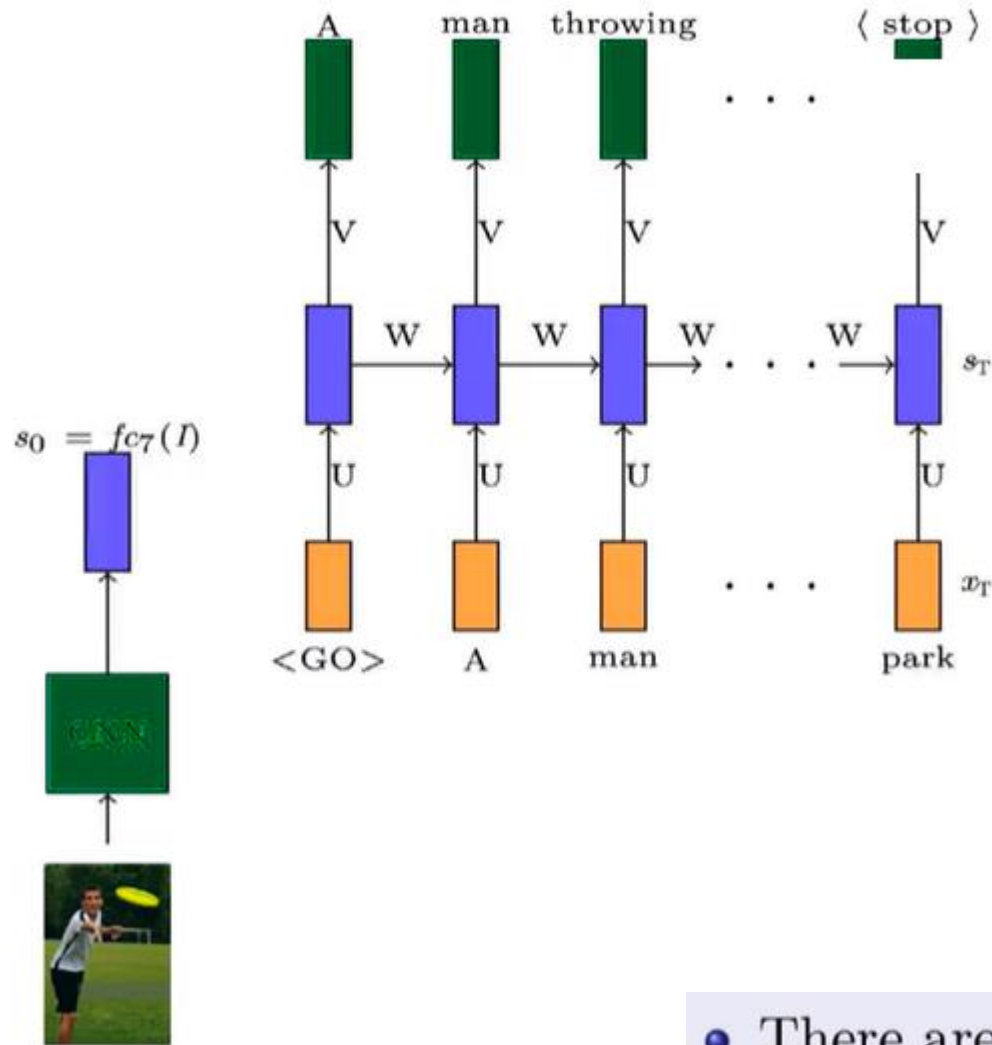
$$P(y_t|y_1^{t-1}) = P(y_t = j|s_t)$$

- Where s_t was a state capturing all the previous words
- We could now model $P(y_t = j|y_1^{t-1}, I)$ as $P(y_t = j|s_t, f_{c7}(I))$



A man throwing
a frisbee in a park

Encoder – Decoder for Image Captioning



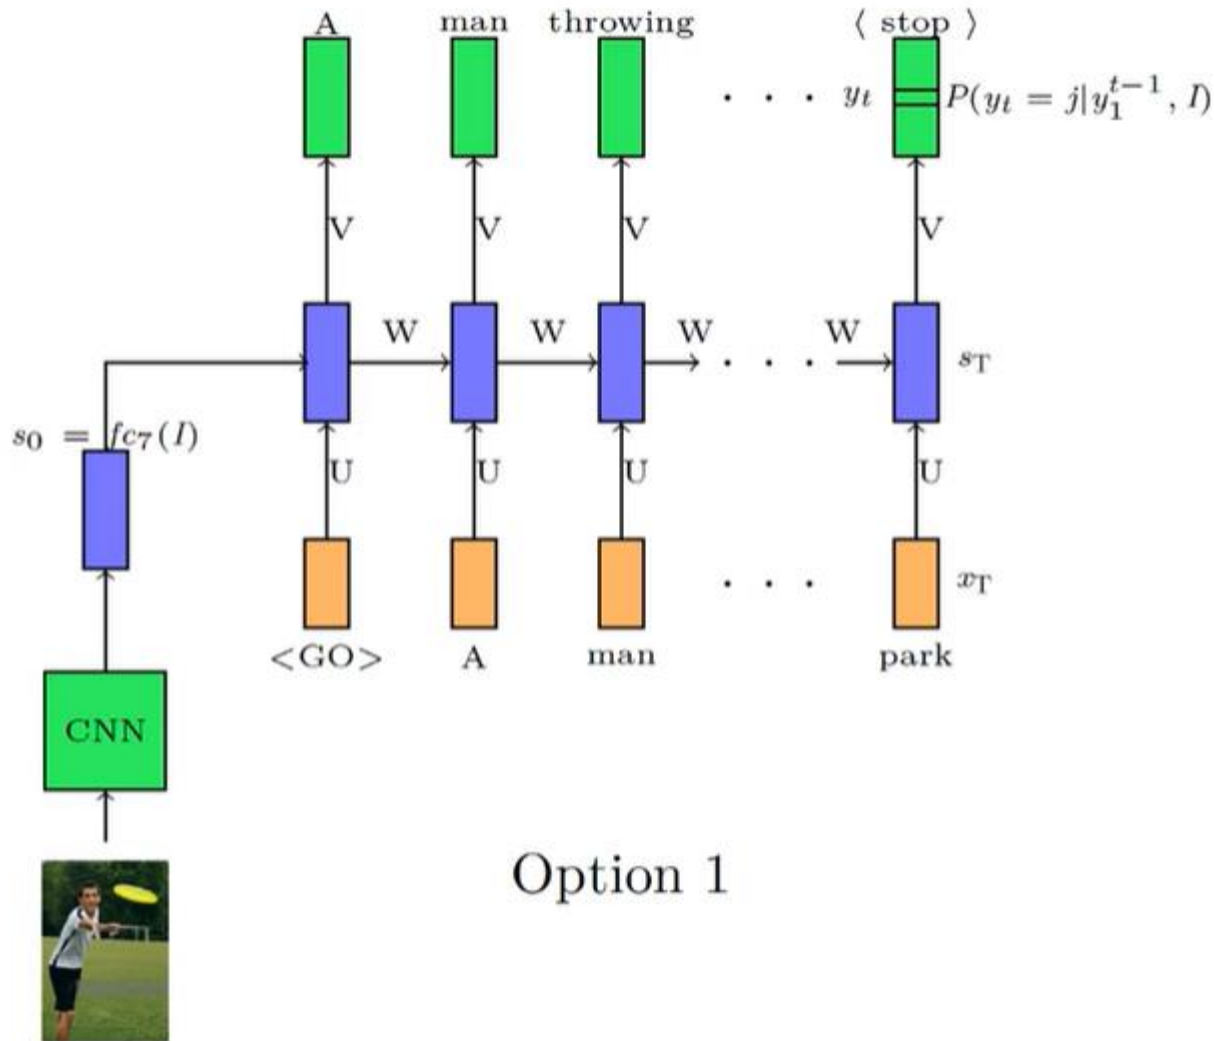
- Earlier we modeled $P(y_t|y_1^{t-1})$ as

$$P(y_t|y_1^{t-1}) = P(y_t = j|s_t)$$

- Where s_t was a state capturing all the previous words
- We could now model $P(y_t = j|y_1^{t-1}, I)$ as $P(y_t = j|s_t, f_{c7}(I))$

- There are many ways of making $P(y_t = j)$ conditional on $f_{c7}(I)$

Encoder – Decoder for Image Captioning



- **Option 1:** Set $s_0 = f_{c7}(I)$
- Now s_0 and hence all subsequent s_t 's depend on $f_{c7}(I)$

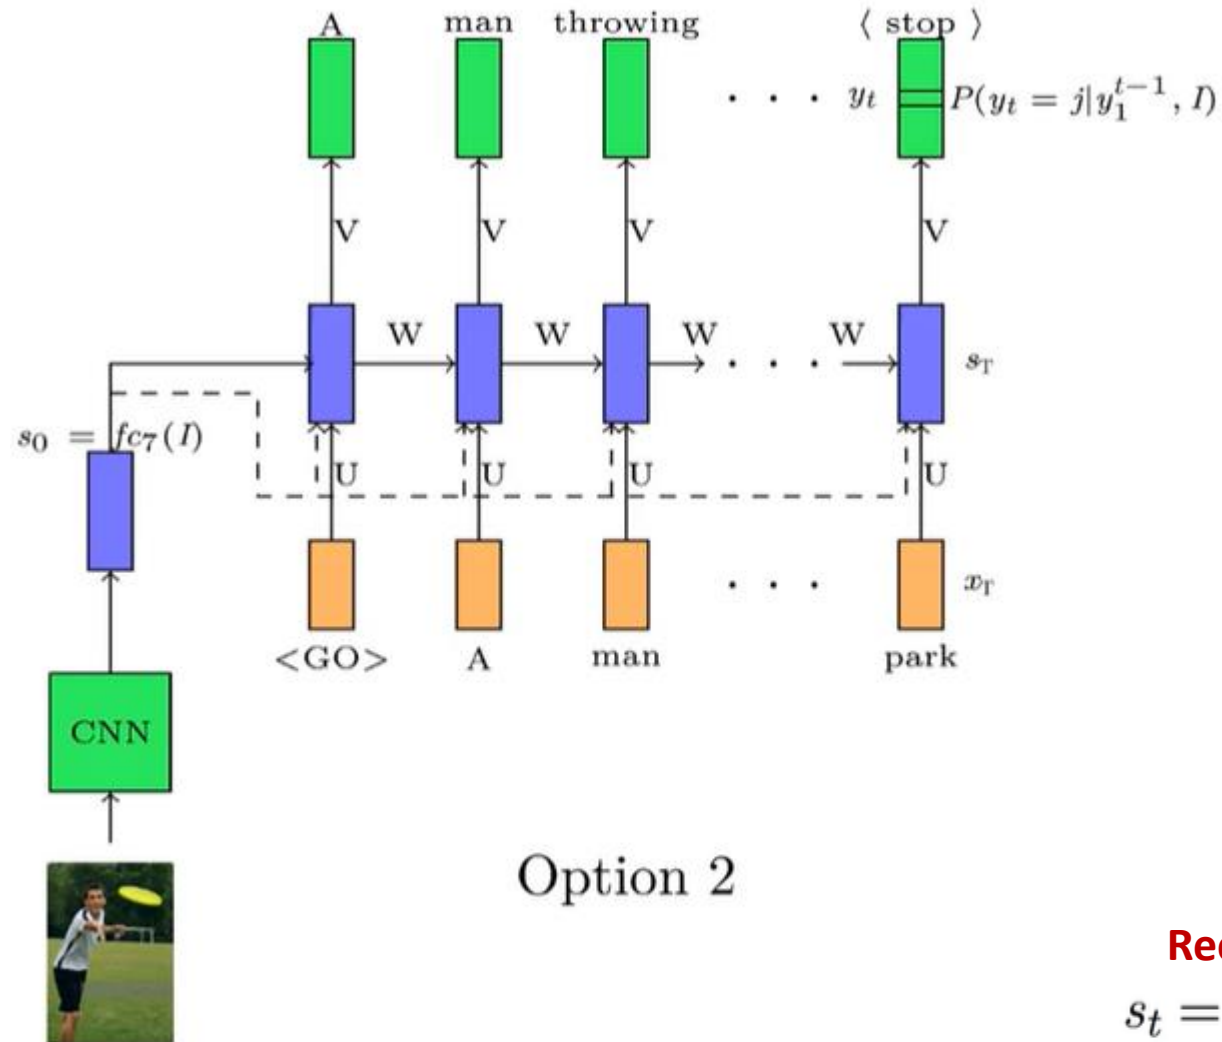
$$S_1 = \sigma(W_{S0} + U_{x1} + b)$$

S_0 computed from CNN, $x_1 \langle GO \rangle$

$$S_2 = \sigma(W_{S1} + U_{x2} + b)$$

Option 1

Encoder – Decoder for Image Captioning



Option 2

- **Option 2:** Another more explicit way of doing this is to compute

$$s_t = \text{RNN}(s_{t-1}, [x_t, f_{c7}(I)])$$

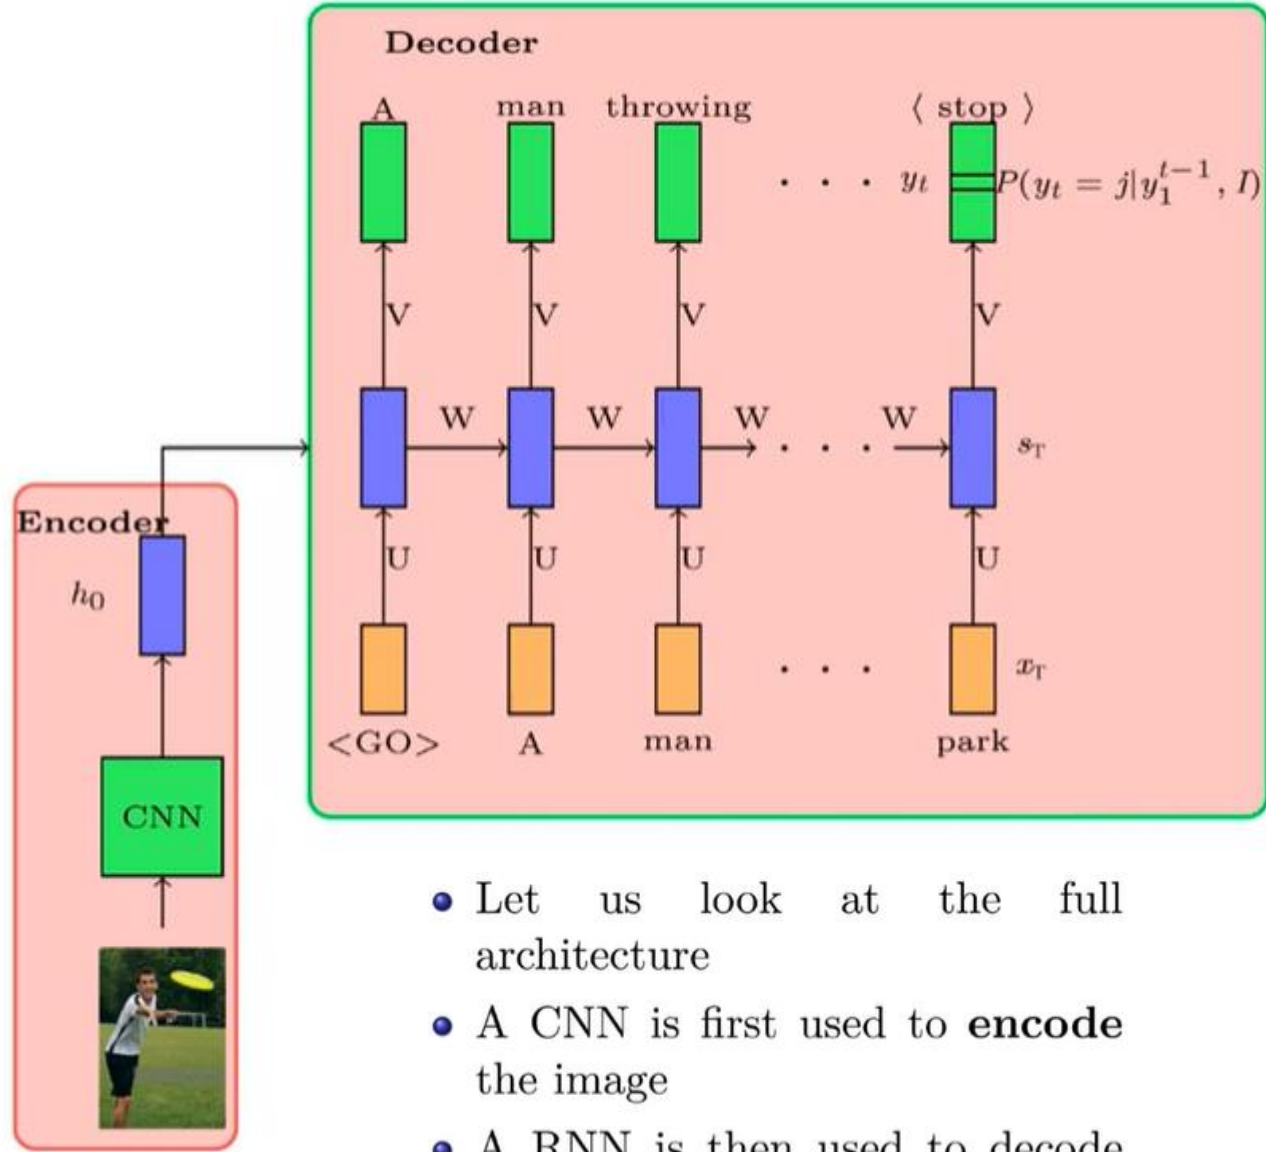
- In other words we are explicitly using $f_{c7}(I)$ to compute s_t and hence $P(y_t = j)$

$$s_2 = \sigma(W_{s1} + U_{[x,1s1]} + b)$$

Concatenate x_1 and s_1

Recap

$$s_t = \sigma(U x_t + W s_{t-1} + b) \quad s_t = \text{RNN}(s_{t-1}, x_t)$$



- Let us look at the full architecture
- A CNN is first used to **encode** the image
- A RNN is then used to decode (generate) a sentence from the encoding

- **Task:** Image captioning
- **Data:** $\{x_i = image_i, y_i = caption_i\}_{i=1}^N$
- **Model:**

- **Encoder:**

$$s_0^* = CNN(x_i)$$

- **Decoder:**

$$s_t = RNN(s_{t-1}, e(\hat{y}_{t-1}))$$

$$P(y_t | y_1^{t-1}, I) = softmax(Vs_t + b)$$


- **Parameters:** $U_{dec}, V, W_{dec}, W_{conv}, b$

- **Loss:**

$$\mathcal{L}(\theta) = \sum_{i=1}^T \mathcal{L}_t(\theta) = - \sum_{t=1}^T \log P(y_t = \ell_t | y_1^{t-1}, I)$$

- **Algorithm:** Gradient descent with backpropagation

Encoder – Decoder for Machine Translation(seq2seq)



- **Task:** Machine translation
- **Data:** $\{x_i = source_i, y_i = target_i\}_{i=1}^N$
- **Model (Option 1):**

- **Encoder:**

$$h_t = RNN(h_{t-1}, x_{it})$$

- **Decoder:**

$$s_0 = h_T \quad (T \text{ is length of input})$$

$$s_t = RNN(s_{t-1}, e(\hat{y}_{t-1}))$$

$$P(y_t | y_1^{t-1}, x) = \text{softmax}(Vs_t + b)$$

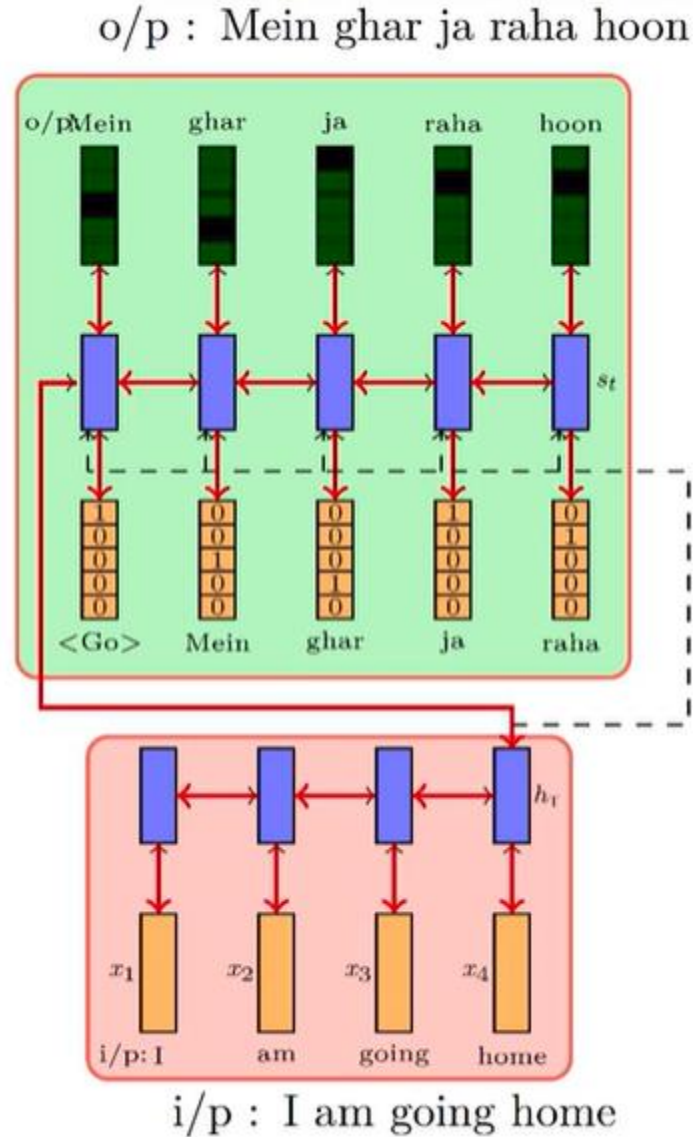
- **Parameters:** $U_{dec}, V, W_{dec}, U_{enc}, W_{enc}, b$

- **Loss:**

$$\mathcal{L}(\theta) = \sum_{i=1}^T \mathcal{L}_t(\theta) = - \sum_{t=1}^T \log P(y_t = \ell_t | y_1^{t-1}, x)$$

- **Algorithm:** Gradient descent with backpropagation

Encoder – Decoder for Machine Translation



- **Task:** Machine translation

- **Data:** $\{x_i = source_i, y_i = target_i\}_{i=1}^N$

- **Model (Option 2):**

- **Encoder:**

$$h_t = RNN(h_{t-1}, x_{it})$$

- **Decoder:**

$$s_0 = h_T \quad (T \text{ is length of input})$$

$$s_t = RNN(s_{t-1}, [h_T, e(\hat{y}_{t-1})])$$

$$P(y_t | y_1^{t-1}, x) = \text{softmax}(Vs_t + b)$$


- **Parameters:** $U_{dec}, V, W_{dec}, U_{enc}, W_{enc}, b$

- **Loss:**

$$\mathcal{L}(\theta) = \sum_{i=1}^T \mathcal{L}_t(\theta) = - \sum_{t=1}^T \log P(y_t = \ell_t | y_1^{t-1}, x)$$

- **Algorithm:** Gradient descent with backpropagation

Encoder – Decoder for Machine Transliteration



- **Task:** Transliteration
- **Data:** $\{x_i = srcword_i, y_i = tgtword_i\}_{i=1}^N$
- **Model (Option 1):**

- **Encoder:**

$$h_t = RNN(h_{t-1}, x_{it})$$

- **Decoder:**

$$s_0 = h_T \quad (T \text{ is length of input})$$


$$s_t = RNN(s_{t-1}, e(\hat{y}_{t-1}))$$

$$P(y_t | y_1^{t-1}, x) = \text{softmax}(Vs_t + b)$$

- **Parameters:** $U_{dec}, V, W_{dec}, U_{enc}, W_{enc}, b$
- **Loss:**

$$\mathcal{L}(\theta) = \sum_{i=1}^T \mathcal{L}_t(\theta) = - \sum_{t=1}^T \log P(y_t = \ell_t | y_1^{t-1}, x)$$

Encoder – Decoder for Machine Transliteration



- **Task:** Transliteration

- **Data:** $\{x_i = srcword_i, y_i = tgtword_i\}_{i=1}^N$

- **Model (Option 2):**

- **Encoder:**

$$h_t = RNN(h_{t-1}, x_{it})$$

- **Decoder:**

$$s_0 = h_T \quad (T \text{ is length of input})$$

$$s_t = RNN(s_{t-1}, [e(\hat{y}_{t-1}), h_T])$$

$$P(y_t | y_1^{t-1}, x) = \text{softmax}(Vs_t + b)$$

- **Parameters:** $U_{dec}, V, W_{dec}, U_{enc}, W_{enc}, b$

- **Loss:**

$$\mathcal{L}(\theta) = \sum_{i=1}^T \mathcal{L}_t(\theta) = - \sum_{t=1}^T \log P(y_t = \ell_t | y_1^{t-1}, x)$$

- **Algorithm:** Gradient descent with backpropagation

Attention Mechanism

o/p : I am **going** home

$t_1 : [1 \ 0 \ 0 \ 0 \ 0]$

$t_2 : [0 \ 0 \ 0 \ 0 \ 1]$

$t_3 : [0 \ 0 \ 0.5 \ 0.5 \ 0]$

i/p : Main ghar **ja raha** hoon

o/p : I am going home

$t_1 : [1 \ 0 \ 0 \ 0 \ 0]$

$t_2 : [0 \ 0 \ 0 \ 0 \ 1]$


$t_3 : [0 \ 0 \ 0.5 \ 0.5 \ 0]$

$t_4 : [0 \ 1 \ 0 \ 0 \ 0]$

i/p : Main ghar ja raha hoon


- Humans try to produce each word in the output by focusing only on certain words in the input
- Essentially at each time step we come up with a distribution on the input words
- This distribution tells us how much attention to pay to each input words at each time step
- Ideally, at each time-step we should feed only this relevant information (i.e. encodings of relevant words) to the decoder

Machine Translation with Attention Mechanism



- We could just take a weighted average of the corresponding word representations and feed it to the decoder
- For example at timestep 3, we can just take a weighted average of the representations of 'ja' and 'raha'
- Intuitively this should work better because we are not overloading the decoder with irrelevant information (about words that do not matter at this time step)
- How do we convert this intuition into a model ?

Machine Translation with Attention Mechanism



- Of course in practice we will not have this oracle
- The machine will have to learn this from the data
- To enable this we define a function


$$e_{jt} = f_{ATT}(s_{t-1}, \mathbf{x}_j)$$

- This quantity captures the importance of the j^{th} input word for decoding the t^{th} output word (we will see the exact form of f_{ATT} later)
- We can normalize these weights by using the softmax function

$$\alpha_{jt} = \frac{\exp(e_{jt})}{\sum_{j=1}^M \exp(e_{jt})}$$

◀ ◻ ▶ ◀ ≡ ▶ ◀ ≡ ▶


Machine Translation with Attention Mechanism



$$\alpha_{jt} = \frac{\exp(e_{jt})}{\sum_{j=1}^M \exp(e_{jt})} \quad c_t = \sum_{j=1}^T \alpha_{jt} h_j$$

- α_{jt} denotes the probability of focusing on the j^{th} word to produce the t^{th} output word
- We are now trying to learn the α 's instead of an oracle informing us about the α 's

Machine Translation with Attention Mechanism




- From now on we will refer to the decoder RNN's state at the t -th timestep as s_t and the encoder RNN's state at the j -th time step as h_j
- Given these new notations, one (among many) possible choice for f_{ATT} is

$$e_{jt} = V_{att}^T \tanh(U_{att}s_{t-1} + W_{att}h_j)$$

- $V_{att} \in \mathbb{R}^d$, $U_{att} \in \mathbb{R}^{d \times d}$, $W_{att} \in \mathbb{R}^{d \times d}$ are additional parameters of the model
- These parameters will be learned along with the other parameters of the encoder and decoder

Machine Translation with Attention Mechanism



- **Task:** Machine Translation

- **Data:** $\{x_i = source_i, y_i = target_i\}_{i=1}^N$

- **Encoder:**

$$h_t = RNN(h_{t-1}, x_t)$$

$$s_0 = h_T$$

- **Decoder:**

$$e_{jt} = V_{attn}^T \tanh(U_{attn} h_j + W_{attn} s_{t-1})$$

$$\alpha_{jt} = \text{softmax}(e_{jt})$$

$$c_t = \sum_{j=1}^T \alpha_{jt} h_j$$

$$s_t = RNN(s_{t-1}, [e(\hat{y}_{t-1}), c_t])$$

$$\ell_t = \text{softmax}(Vs_t + b)$$

- **Parameters:** $U_{dec}, V, W_{dec}, U_{enc}, W_{enc}, b, U_{attn}, V_{attn}$

- **Loss and Algorithm** remains same

Namah Shivaya